



# Machine Learning Internship at Oculo

## Inertial Odometry via IMU-only state estimation

**John-Paul Marletta**

Sup. Robert Kesten - Technical Lead

MSci Physics, Royal Holloway, University of London

Technical Presentation  
October 16, 2023

# Presentation Structure

- 1 Project Overview
- 2 Literature Review
- 3 Data formatting for TLIO
- 4 DevOps / Cloud Computing
- 5 Training TLIO
- 6 Results and Analysis
- 7 References

# Project Overview

- **Position:** ML Engineer (Intern)
- **Project Title:** 'Inertial Odometry and trajectory estimation'
- **Goals**
  - Data formatting, data pipelining, dataset for compatibility with chosen model
  - To train and test a **CNN** [Chen et al], , for IMU-only state estimation
  - This is to improve the overall tracking accuracy of Oculo's product by **outputting trajectory estimation** when the SLAM algorithm fails.

# Literature Review

Promising case: IONet [Chen et al.]

## Expectation: completely independent:



- Conducted my own LR
- Proposed papers with most promising models.

## Promising model: IONet

- Approaches prior to IONet:
  - **SINS**: direct IMU data integration → **error prone**,
  - **PDR**: integration + step detection algorithms
- IONet **break continuous integration**, instead segment inertial data into semi-independent windows

# Literature Review

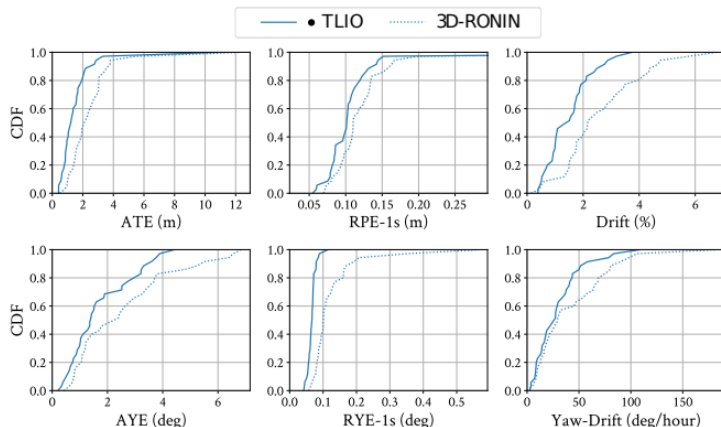
Chosen model: TLIO - Tight Learned Inertial Odometry [Liu et al.]



- UPenn & **Meta** labs - 2020
- Builds on IONet, now regressing **3D** trajectories
- Propose complete state-estimation system combining **CNN/ResNet**, IMU **Buffer**, with **EKF**
- **EKF** estimates position, orientation, velocity, and IMU biases with only pedestrian IMU data
- **Reduces average yaw and position drift by 27% and 33%**, compared to best RoNIN 2019 model, using velocity concatenation approach.

# Literature Review

Chosen model: TLIO's performance



# Data formatting for TLIO

Overview and skills learnt



SciPy



pandas

- **Numpy, SciPy, pandas, matplotlib**
- Developed **File I/O** skills : iterated over large dataset
- Approx. **700 line** script
- Greatly improved **de-bugging**, also via Run and Debug and learnt **Unit-Testing**.
- **Data pipeline:** OxIOD → **Parsing** csv files → interpolating/formatting → dicts → converting to **pandas** data frames → csv files → hdf5 generation script → hdf5 files

# Data formatting for TLIO

## Interpolating data and Quaternions

- Wrote function for ground-truth (GT) **velocity estimations** from motion capture positions and timestamps.
- **Linear interpolation** (`scipy.interpolate.interp1d`) of motion capture data aligning with IMU-only data timestamps.
- **Spherical linear interpolation of quaternions**(`pyquaternion.Quaternion.slerp`) representing rotations - more computationally efficient.



# Data Formatting for TLIO

## Example of Interpolating Quaternions

```
def _interpolate_ox_vicon_quaternions(ox_imu_upsampled_data, quaternions, clipped_ox_vicon_data):
    """ - Interpolating vicon quaternions using speherical linear interpolation
        - creating dictionary: ox_interpolated_vicon_data and adding components from interpolated_quaternions
        list into separate lists inside ox_interpolated_vicon_data
    """
    ox_interpolated_vicon_data = {'timestamps': ox_imu_upsampled_data['timestamps'],
                                  'rotation_w': [], 'rotation_x': [], 'rotation_y': [], 'rotation_z': [],
                                  'translation_x': [], 'translation_y': [], 'translation_z': []}

    vicon_clipped_timestamps = clipped_ox_vicon_data['timestamps']

    for timestamp in ox_imu_upsampled_data['timestamps']:
        original_index = bisect.bisect(vicon_clipped_timestamps, timestamp) - 1

        if original_index >= len(vicon_clipped_timestamps) - 1:
            original_index = len(vicon_clipped_timestamps) - 2

        if original_index < 0:
            original_index = 0

        fraction = ((timestamp - vicon_clipped_timestamps[original_index]) /
                    (vicon_clipped_timestamps[original_index + 1] - vicon_clipped_timestamps[original_index]))

        slerped_quat = Quaternion.slerp(quaternions[original_index], quaternions[original_index + 1], fraction)

        ox_interpolated_vicon_data['rotation_w'].append(slerped_quat.w)
        ox_interpolated_vicon_data['rotation_x'].append(slerped_quat.x)
        ox_interpolated_vicon_data['rotation_y'].append(slerped_quat.y)
        ox_interpolated_vicon_data['rotation_z'].append(slerped_quat.z)

    print("quaternions interpolated")

    return ox_interpolated_vicon_data
```

## Docker

- Lightweight for running TLIO remotely
- Building docker containers, images, mounting volumes
- Implemented '**keepalive**' feature in **ssh config file** to solve connection issues while copying TLIO to ec2 as **.tar.gz**
- **Adjusted container memory**, fixed issue
- Used docker container with **conda** to solve local system issue



## AWS - EC2

- Ran docker in EC2 VM, via ssh, for CPU and memory strain
- Used **r5.2xlarge** machine: **8 vCPUs** for **multiprocessing**, 64GB
- **CLI: git bash**
- **Linux commands** including file operations and **cat**, **scp**, **ssh**, **ping**, **logs**, **top**, **ps** etc..

# Training TLIO

- **TLIO input dim:**  $N \times 6$ , ( $N$  IMU samples) - gravity-aligned
- **Multiprocessing** - 7 workers
- Adam optimisation
- **80%, 10%, 10%** train, test, validate split
- batch sizes of 1024

$$L_{\text{MSE}}(\mathbf{d}, \hat{\mathbf{d}}) = \frac{1}{n} \sum_{i=1}^n \left\| \mathbf{d}_i - \hat{\mathbf{d}}_i \right\|^2$$

- $\mathbf{d}$  = 3D output displacement from model,  $\hat{\mathbf{d}}$  is ground truth data,  $n$  is the number of data points in the training set

# Results and Analysis

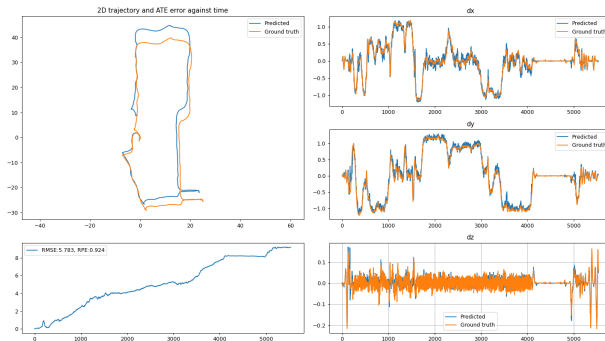



Figure: 10th training epoch: estimated vs ground truth trajectories


- **Drift:** within 0.02 – 0.05 in m/m
- **MSE loss:**  $x$  : (0.006 – 0.152),  $y$  : (0.0004 – 0.13896),  $z$  : (0.00212 – 0.17467)
- More training was needed at this point.

# Results and Analysis



**Figure:** Plots of  $x$ ,  $y$ ,  $z$  predictions vs Ground Truth, over trajectory, at 10th training epoch

## Thank you! Questions? Comments?

 Liu, Wenxin, et al. (2020)  
TLIO: Tight learned inertial odometry.  
*Robotics and Automation Letters* 5.4: 5653-5660.

 Chen, Changhao, et al (2018)  
IONet: Learning to cure the curse of drift in inertial odometry.  
*Proceedings of the AAAI Conference on Artificial Intelligence*. Vol. 32. No. 1.

 Liue, Wenxin, et al. (2020)  
<https://github.com/CathIAS/TLIO>